

Setup Files

Data Check

```
[hdfs@big-nn01 /]$ cd /training/analyst/data/Sample\ Examination/
```

Import Test Files

```
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -mkdir /user/exam
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem1 /user/exam/Problem1
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem2 /user/exam/Problem2
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem3 /user/exam/Problem3
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem4 /user/exam/Problem4
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem5 /user/exam/Problem5
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem6 /user/exam/Problem6
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem7 /user/exam/Problem7
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem8 /user/exam/Problem8
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem9 /user/exam/Problem9
[hdfs@big-nn01 Sample Examination]$ hdfs dfs -put Problem10 /user/exam/Problem10
```

Problem 1

Instructions

Write a query to compare each active account's balance to the average balance of all active accounts of the same type.

Data Description

The account data exists in the metastore account table in the problem1 database. The amount column gives the account balance. The type column gives the type of account and the status column gives the status of the account.

Output Requirements

1. Write the report query in the local file: /home/training/problem1/solution.sql
2. Executing the solution.sql script from the hive command-line tools should generate the report output
3. Only include accounts for which status is Active
4. Create a difference column holding the difference between the account balance (amount) and the average balance for accounts of the same type to the right of all the other columns
5. The query should return the following columns from the account table: id, type, status, amount, as well as the difference calculated

-
- Create Database

```
create database exam;  
use exam;
```

- Make Table in Hive

```
create external table account (  
id STRING, status STRING, amount INT, type STRING)  
ROW FORMAT DELIMITED  
FIELDS TERMINATED BY ','  
LOCATION '/user/exam/Problem1'
```

- Excute Query

```
select a.id, a.type, a.status, a.amount,  
(b.average - a.amount) as difference  
from account a  
join  
(select type, avg(amount) as average from account where status = "Active"  
group by type) b  
on (a.type = b.type)  
where status = "Active";
```

3 / 20

Problem 2

Instructions

Create an employee table in the metastore that contains the employee records stored in HDFS.

Data Description

All of the employee records are stored in the /user/training/problem2/data/employee/ HDFS directory in Parquet file format. The files contain the following columns and types:

Output Requirements

1. The table you create should be named solution and stored in the problem2 database
2. The table must point to the existing data in HDFS

- Make Table

```
CREATE EXTERNAL TABLE solution (
  id int, fname string, lname string, address string, city string, state string, zip string, birthday string,
  hireday string )
stored as parquet
location "/user/exam/Problem2/data/employee/"
```

The screenshot shows the Hive web interface. On the left, a sidebar displays the database structure under 'exam', including a table named 'solution' with columns: id (int), fname (string), lname (string), address (string), city (string), state (string), zip (string), birthday (string), and hireday (string). The main area shows a SQL query being executed:

```
1 CREATE EXTERNAL TABLE solution (
2 id int,
3 fname string,
4 lname string,
5 address string,
6 city string,
7 state string,
8 zip string,
9 birthday string,
10 hireday string )
11 stored as parquet
12 location "/user/exam/Problem2/data/employee/"
```

Below the query, the execution log shows the following output:

```
location /user/exam/Problem2/data/employee/
INFO : Starting task [Stage-0:DDL] in serial mode
INFO : Completed executing command(queryId=hive_20190613014545_b6b6c50f-5ad2-4893-86f8-19c40350674e);
Time taken: 0.059 seconds
INFO : OK
```

Problem 3

Instructions

Generate a table that contains all customers who have negative account balances.

Data Description

The customer records are stored in the customer table in the problem3 database. The account records are stored in the account table in the problem3 database. The account records contain a field called amount that may be negative. The custid field is a foreign key into the customer table.

Output Requirements

1. Create a metastore table named solution stored in the problem3 database that contains the customer records that match the criteria
2. Each solution record should contain the Customer ID, first name, last name, and home phone of the customer
3. Maintain the same column names and datatypes in the new table as the fields from the customer table

- Make Table

```
create external table solution3
(id STRING, fname string, lname string, hphone string)
ROW FORMAT DELIMITED
FIELDS TERMINATED BY ','
LOCATION '/user/exam/Problem3/customer'
```

The screenshot shows the Hue web interface. On the left, the 'exam' database is selected, and the 'solution3' table is listed. The main area shows the query editor with the query 'SELECT * FROM solution3;'. Below the query editor, the execution log shows the command being executed and the results. The results are displayed in a table with 6 rows and 4 columns: solution3.id, solution3.fname, solution3.lname, and solution3.hphone.

	solution3.id	solution3.fname	solution3.lname	solution3.hphone
1	10000	Cullen	Johnson	(504) 159-1501
2	10001	Sybil	Wiley	(504) 780-0366
3	10002	Victoria	Mcpherson	(990) 307-7932
4	10003	Deanna	Pollard	(178) 455-7090
5	10004	Dominic	Aguilar	(647) 503-9092
6	10005	Jana	Patrick	(751) 459-6046

Problem 4

Instructions

LoudAcre Mobile has merged with another company located in California. Each company has a list of customers in different formats. Combine the two customer lists into a single dataset using an identical schema.

Data Description

The original customer data exists in the HDFS directory `/user/training/problem4/data/employee1/`. It contains expanded, nine-digit zip codes. The new files are in the HDFS directory `/user/training/problem4/data/employee2/`. It contains last names before first names, both using all capital letters.

Output Requirements

1. Combine these files into a single tab-delimited dataset and stored in the HDFS directory `/user/training/problem4/solution/`
2. Each record should be in the following format:
3. Only include customers whose state is 'CA'

Excute on PIG

```
employee1 = LOAD '/user/training/problem4/data/employee1' AS (customerID:Int, fname:chararray, lname:chararray, address:chararray, city:chararray, state:chararray, zip: chararray);
```

```
e1 = FOREACH employee1 GENERATE customerID, fname, lname, address, city, state, SUBSTRING(zip,0,5);
```

```
employee2 = LOAD '/user/training/problem4/data/employee2' USING PigStorage(',') AS (customerID:Int, junk:Int, lname:chararray, fname:chararray, address:chararray, city:chararray, state:chararray, zip: chararray);
```

```
e2 = FOREACH employee2 GENERATE customerID, UCFIRST(LOWER(fname)), UCFIRST(LOWER(lname)), address, city, state, zip;
```

```
both = UNION e1 , e2;
```

```
STORE both INTO '/user/training/problem4/solution/'
```

UE

Query Editors

Data Browsers

Workflows

Security

File Browser

ACTIONS

View as binary

Edit file

Download

View file location

Refresh

INFO

Last modified

June 12, 2019 6:19 p.m.

User

training

Group

supergroup

Size

6.3 KB

Mode

Home

/ user / training / problem4 / solution / part-m-00000

Page 1 of 2

10000000	Oiga	Booker	Ap #643-2741 Proin Street	Gresham OR	42593		
10000001	Raja	Spence	P.O. Box 765, 7700 Eros Rd.	Duluth MN	67110		
10000002	Meredith	Schwartz	3414 At Road	San Antonio TX	35713		
10000003	Thor	Lloyd	2703 Amet, Road Rock Springs	WY 73861			
10000004	Arden	Mooney	Ap #888-1187 Aliquam Road	Rockville MD	82478		
10000005	Kenneth	Lucas	Ap #637-8746 Feugiat Av.	Boston MA	84047		
10000006	Leonard	Goodwin	Ap #465-7707 Aliquam Avenue	Rock Springs WY	91337		
10000007	Elton	Conrad	P.O. Box 787, 8889 Nam Street	Auburn ME	28280		
10000008	Ryan	Keith	6857 Et Ave Springfield	MA 58872			
10000009	Griffin	Combs	Ap #266-8677 Nulla Street	Ketchikan AK	99904		
10000010	Jasmine	Ramos	P.O. Box 486, 5722 Faucibus Av.	Tallahassee FL	34598		
10000011	Dolan	King	P.O. Box 382, 542 Mauris Av.	Grand Rapids MI	93708		
10000012	Barrett	Serrano	P.O. Box 143, 2159 Ac, Road	Memphis TN	83389		
10000013	Maggy	Mcdaniel	7096 Enim Avenue	Davenport IA	73580		
10000014	Erich	George	8539 Ut Ave Owensboro	KY 56812			
10000015	Kirestin	Rowland	P.O. Box 764, 1149 Consequat Rd.	West Valley City UT	24551		
10000016	Phoebe	Bowen	932-4942 Laoreet Road	Toledo OH	47662		
10000017	Myles	Oneill	Ap #376-6356 Amet, St.	Huntsville AL	35832		
10000018	Thaddeus	Vinson	2787 Orci, Ave	Huntsville AL	36797		

8 / 20

Problem 5

Instructions

The bank is making a Facebook group for the Palo Alto, CA branch. Generate a script that outputs the customers and employees who live in Palo Alto, CA.

Data Description

The employee records are stored in the employee metastore table in the problem5 database, The customer records are stored in the customer metastore table in the problem5 database.

Output Requirements

1. Write the report query in the local file /home/training/problem5/solution.sql
2. Executing the solution.sql script from the hive command-line tools should generate the report output
3. The report should contain first name, last name, city, and state with a tab as the delimiter between fields
4. Only output records that have a city value of 'Palo Alto' and state value of 'CA'
5. Duplicate records (if any) should be included

-
- Make Table

```
create external table customer5 ( id int, fname string, lname string, address string, city string, state string, zip string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LOCATION "/user/exam/Problem5/customer";
```

```
create external table employee5 ( id int, fname string, lname string, address string, city string, state string, zip string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LOCATION "/user/exam/Problem5/employee";
```

- Excute Query

```
select fname, lname, city, state from customer5 where (city = "Palo Alto" AND state = "CA") UNION ALL select fname, lname, city, state from employee5 where (city = "Palo Alto" AND state = "CA");
```

10 / 20

Problem 6

Instructions

There are privacy concerns about the employee data that is stored on the cluster. Your task is to remove any age information from the employee data by creating a new table for the data analysts to query against.

Data Description

All of the employee records are stored in the employee metastore table in the problem6 database.

Output Requirements

1. Create a new table named solution stored in the problem6 database with the same file format as the employee table
2. Maintain the same column names and datatypes in the new table
3. The birthday field in the solution table should be truncated to only contain month/day instead of the current month/day/year data that is in the employee table

```
create external table employee6 ( id int, fname string, lname string, address string, city string, state string, zip string, birthday string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LOCATION "/user/exam/Problem6/employee";
```

```
create table solution6 as select id, fname, lname, address, city, state, zip, substr(birthday,1,5) from employee6;
```

The screenshot displays the Hue web interface. On the left, a sidebar shows the database structure with 'exam' selected, listing tables like 'account', 'customer5', 'employee', 'employee5', 'employee6', 'solution', 'solution3', and 'solution6'. The main panel shows a Hive query execution log with the following output:

```
FROM solution6
INFO : Completed executing command(queryId=hive_20190613062020_5ff88177-d732-443b-93c1-c7ab5fbffaeb);
Time taken: 0.002 seconds
INFO : OK
```

Below the log, the query results are displayed as a table with 14 rows and 4 columns: solution6.id, solution6.fname, solution6.lname, and solution6.address.

	solution6.id	solution6.fname	solution6.lname	solution6.address
1	10000000	Deanna	Lane	900-1514 Vitae, Rd.
2	10000001	Hall	Garrett	9656 Urna Avenue
3	10000002	Lucian	Dotson	P.O. Box 277, 4808 Fusce St.
4	10000003	Yuri	Sherman	Ap #399-8275 Molestie Road
5	10000004	Jaime	Griffin	Ap #647-2123 Quis Rd.
6	10000005	Zorita	Weber	747-9424 Orci, Av.
7	10000006	Mara	Meadows	517-4594 Ac, Rd.
8	10000007	Evan	Richard	P.O. Box 223, 8182 Non, Av.
9	10000008	Briar	Anderson	Ap #548-6452 Nunc Road
10	10000009	Cole	Odom	P.O. Box 962, 2496 Sodales St.
11	10000010	Edward	Dorsey	8473 Aliquam Rd.
12	10000011	Amery	Mason	Ap #231-8154 Dictum Rd.
13	10000012	Keaton	Warren	Ap #173-7618 Tellus Rd.
14	10000013	Erin	Ford	112-1682 Mauris Rd.

Problem 7

Instructions

Generate a report that contains all of the Seattle employee names in sorted order.

Data Description

The employee records are stored in the employee table in the problem7 database.

Output Requirements

1. Write the report query in the local file /home/training/problem7/solution.sql
2. Executing the solution.sql script from the hive command-line tools should generate the report output
3. The employee names should be printed out as the first name, a space, and the last name
4. The output should be sorted by first name and then by last name and should only contain employees whose city is 'Seattle'
5. Duplicate names should be included (if any)

-
- Make Table

```
create external table employee7 ( id int, fname string, lname string, address string, city string, state string, zip string, birthday string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LOCATION "/user/exam/Problem7/employee";
```

- Excute Query

```
select concat(fname, " ", lname) as name from employee7 where city = "Seattle" order by name;
```

쿼리 기록 저장된 쿼리 결과 (3)

	name
1	Lucian Dotson
2	Zeph Horn
3	Zeph McClain

Problem 8

Instructions

Use Sqoop to export customer data from HDFS into a MySQL database table. Place the data in the solution table in MySQL, which has been created and is currently empty.

Data Description

The data files are in the HDFS directory /user/training/problem8/data/customer/. MySQL database information:

1. Installation: localhost
2. Database name: problem8 | Table name: solution
3. Username: cloudera
4. Password: cloudera

Output Requirements

1. Export all of the customer data from HDFS into the MySQL solution table
2. The solution table already exists in the MySQL database but currently has no rows

```
sqoop export
--connect jdbc:mysql://localhost/problem8
--username cloudera
--password cloudera
--table solution
--input-fields-terminated-by '\t'
--export-dir hdfs://localhost:8020/user/training/problem8/data/customer
```

```
mysql> select * from solution;
```

id	fname	lname	address	city	state	zip	birthday
10000000	Deanna	Lane	900-1514 Vitae, Rd.	Lafayette	LA	97827	08/31/2016
10000001	Hall	Garrett	9656 Urna Avenue	Tucson	AZ	86511	08/24/2016
10000002	Lucian	Dotson	P.O. Box 277, 4808 Fusce St.	Seattle	WA	57731	08/12/2016
10000003	Yuri	Sherman	Ap #399-8275 Molestie Road	Kapolei	HI	16943	08/26/2016
10000004	Jaime	Griffin	Ap #647-2123 Quis Rd.	Madison	WI	51394	08/13/2016
10000005	Zorita	Weber	747-9424 Orci, Av.	Hattiesburg	MS	90262	08/09/2016
10000006	Mara	Meadows	517-4594 Ac, Rd.	Huntsville	AL	35374	08/11/2016
10000007	Evan	Richard	P.O. Box 223, 8182 Non, Av.	College	AK	99682	08/25/2016
10000008	Briar	Anderson	Ap #548-6452 Nunc Road	Cleveland	OH	90704	08/18/2016
10000009	Cole	Odom	P.O. Box 962, 2496 Sodales St.	Boston	MA	27282	08/21/2016
10000010	Edward	Dorsey	8473 Aliquam Rd.	Montgomery	AL	36204	08/13/2016
10000011	Amery	Mason	Ap #231-8154 Dictum Rd.	Memphis	TN	58035	08/09/2016
10000012	Keaton	Warren	Ap #173-7618 Tellus Rd.	Tallahassee	FL	19431	08/20/2016
10000013	Erin	Ford	112-1682 Mauris Rd.	Nashville	TN	51731	08/26/2016
10000014	Madison	Oconnor	2326 Velit Rd.	Dover	DE	93480	08/19/2016
10000015	Baker	Hodges	P.O. Box 932, 4163 Cursus Rd.	Salem	OR	16709	08/25/2016
10000016	Leo	Kane	Ap #832-2102 Malesuada Rd.	Glendale	AZ	86065	08/14/2016
10000017	Kyra	Hays	P.O. Box 311, 5471 Faucibus Road	Springfield	IL	10229	08/21/2016
10000018	Wallace	Saunders	Ap #462-9715 Massa. Road	Nampa	ID	46348	08/16/2016
10000019	Anjolie	Whitaker	7004 Dis St.	Bloomington	MN	69004	08/13/2016
10000020	Price	Washington	911-5979 Non, Ave	Bloomington	MN	14550	08/11/2016
10000021	Zeph	Mcclain	Ap #248-9912 Elementum, St.	Seattle	WA	52142	08/25/2016
10000022	Sopoline	Hall	P.O. Box 860, 8554 Sed Street	Provo	UT	28779	08/20/2016
10000023	Riley	Mcfadden	539-6573 Vitae, Avenue	Toledo	OH	24078	08/08/2016
10000024	Kareem	Houston	P.O. Box 522, 594 Ligula. St.	Tucson	AZ	86555	08/12/2016
10000025	Fatima	Rowe	515-1899 Duis St.	Oklahoma City	OK	66663	08/21/2016
10000026	Fallon	Middleton	4272 Sed Ave	Fort Smith	AR	71864	08/27/2016

Problem 9

Instructions

Your company is being acquired by another company. To prepare for this acquisition, update the customer records to guarantee there will be no duplicate IDs with their existing customer IDs.

Data Description

The customer records are stored in the customer table in the problem9 database. The id column is a unique identifier for that record.

Output Requirements

1. Create a new table named solution stored in the problem9 database
2. Maintain the same column names and datatypes as the customer table, except store the id as a string
3. The solution table should have all of the data from the customer table, with the addition of the letter 'A' to the existing id values to make them unique

-
- Make Table

```
create external table customer9 ( id int, fname string, lname string, address string, city string, state string, zip string, birthday string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t' LOCATION "/user/exam/Problem9/customer";
```

- Excute Query

```
create table solution9 as select concat("A", cast(id as string)) as id, fname, lname, address, city, state, zip, birthday from customer9;
```

[illegible]

Problem 10

Instructions

Your boss needs specialized reports using the billing data and is constantly asking for help to write SQL queries. Create a database view in the metastore so that your boss has customer and billing data joined.

Data Description

The customer data exists in the customer metastore table in the problem10 database. The billing data exists in the billing metastore table in the problem10 database. The id column is a foreign key into which customer has the charge.

Output Requirements

1. Create a new view named solution stored in the problem10 database
2. The new view should maintain the same column datatypes as the source tables
3. Do not return customer records that have no billing data, or billing records that have no matching customer
4. The new view should contain the following columns
5. The billdate column should only contain the date field and no time information

- Make Table

```
create external table customer10 ( id int, fname string, lname string, address string, city string,
state string, zip string, birthday string) ROW FORMAT DELIMITED FIELDS TERMINATED BY '\t'
LOCATION "/user/exam/Problem9/customer";
```

```
create external table billing10 ( id int, charge double, tstamp string) ROW FORMAT DELIMITED
FIELDS TERMINATED BY '\t' LOCATION "/user/exam/Problem10/billing";
```

- Make View

```
create view solution as select c.id as id, c.fname as fname, c.lname as lname, c.city as city, c.state
as state, b.charge as charge, SUBSTR(b.tstamp,0,10) as billdate from customer c join billing b on
(c.id = b.id);
```

Hive

Query Editors

Data Browsers

Workflows

Security

Add a name...

Add a description...

accounting

Tables

(4)

accounts

billing10

customer10

solution10

Query History

Saved Queries

Results

	solution10.id	solution10.fname	solution10.lname	solution10.city	solution10.state	solution10.charge	solution10.billd
1	1000000	Medge	Roach	Racine	WI	15.789999999999999	2017-03-05
2	1000001	Nasim	Stone	Tuscaloosa	AL	57.729999999999997	2016-09-05
3	1000002	Jolie	Schneider	Kapolei	HI	556.03999999999996	2017-02-06
4	1000003	Lacota	Molina	Philadelphia	PA	654.63	2016-12-01
5	1000004	Blaine	Sweet	Topeka	KS	287.12	2017-05-30
6	1000005	Lesley	Bird	Dallas	TX	900.13999999999999	2016-10-26
7	1000006	Sydnee	Howell	Dallas	TX	347.88	2016-10-04
8	1000007	Jermaine	Griffin	Baltimore	MD	754.84000000000003	2017-08-01
9	1000008	Brynn	Pennington	Lincoln	NE	289.66000000000003	2017-03-26
10	1000009	Ava	Noble	Baton Rouge	LA	928.67999999999995	2016-12-12
11	1000010	Armand	Rodriguez	Great Falls	MT	211.03	2017-04-20

training@localhost:~

Problem 11

Instructions

Several analysis questions are described below and you will need to write the SQL code to answer them. You can use whichever tool you prefer – Impala or Hive – using whichever method you like best, including shell, script, or the Hue Query Editor, to run your queries.

Data Description

Using default database from the metastore for these queries.

Output Requirements

1. Write the report query in the local file /home/training/problem11/solution.sql
2. Executing the solution.sql script from the hive command-line tools should generate

the report output

- a. Which top three products has Dualcore sold more of than any other? Hint: Remember that if you use a GROUP BY clause, you must group by all fields listed in the SELECT clause that are not part of an aggregate function.

```
SELECT c.name, count(*) FROM orders a, order_details b, products c WHERE a.order_id =
b.order_id AND b.prod_id = c.prod_id AND c.brand = 'Dualcore' GROUP BY c.name LIMIT 3;
```

The screenshot shows the Hue Query Editor interface. On the left is a sidebar with a 'Tables' section listing database tables: ads, customers, employees, lation, order_details (with columns order_id (int), prod_id (int)), orders (with columns order_id (int), cust_id (int), order_date (timestamp)), products, ratings, and solution. The main area displays a SQL query in a text editor:

```
1 SELECT c.name, count(*)
2 FROM orders a,
3      order_details b,
4      products c
5 WHERE a.order_id = b.order_id
6       AND b.prod_id = c.prod_id
7       AND c.brand = 'Dualcore'
8 GROUP BY c.name
9 LIMIT 3;
```

Below the query editor, the 'Results' tab is active, showing a table with two columns: 'c.name' and '_c1'. The table contains three rows of data:

	c.name	_c1
1	1.5 TB SATA3 Disk	3956
2	16 GB Micro SD	3279
3	2 GB Micro SD	1979

- b. Calculating Revenue and Profit – write a query to show Dualcore's revenue (total price of products sold) and profit (price minus cost) by date. Hint: The order_date column in the orders table is of type TIMESTAMP. Use the function to_date to get just the date portion of the value.

```
SELECT to_date(a.order_date), sum(c.price), sum(c.price - c.cost)
FROM orders a, order_details b, products c WHERE a.order_id = b.order_id AND b.prod_id =
c.prod_id AND c.brand = 'Dualcore' GROUP BY to_date(a.order_date);
```

	_c0	_c1	_c2
1	2008-06-01	170742	14018
2	2008-06-02	270806	26356
3	2008-06-03	157680	16130
4	2008-06-04	84182	9349
5	2008-06-05	99186	10465
6	2008-06-06	142582	11095
7	2008-06-07	217946	25676
8	2008-06-08	93020	12750
9	2008-06-09	64133	7884
10	2008-06-10	127639	14730
11	2008-06-11	100127	15137
12	2008-06-12	175368	15097

- c. Calculating the order Total – Which ten orders had the highest total dollar amounts?

```
SELECT a.order_id, SUM(c.price) as Total FROM orders a, order_details b, products c WHERE
a.order_id = b.order_id AND b.prod_id = c.prod_id GROUP BY a.order_id ORDER BY total desc
LIMIT 10;
```

	a.order_id	total
1	5605465	940577
2	5997571	702157
3	5551963	699348
4	5944419	627266
5	5401363	624428
6	6156005	